

Lecture 26: Parallelism in the Real World, Design of 15-210, and Concurrency

Lecturer: Umut A. Acar

Why Parallelism?

- Desire to solve larger more important problems.
 - One machine/processor insufficient
 - Multiple machines needed.
 - Possible because problems often decompose.
- Examples:
 - Scientific computing: massive computations
 - Internet: massive data, many queries
 - Big data

Cray-1 (1976): the world's most expensive love seat



Data Center: Hundred's of thousands of computers



Since 2005: Multicore computers



AMD Opteron (sixteen-core) Model 6274

by [AMD](#)

(1 customer review)

List Price: \$693.00

Price: \$599.99 Prime

You Save: \$93.01 (13%)


Only 1 left in stock (more on the way).

Ships from and sold by [Amazon.com](#). Gift-wrap available.

Want it delivered Monday, November 5? Order it in the next 14 hours and 37 minutes


Delivery may be impacted by Hurricane Sandy. Proceed to checkout to see estimated

[43 new from \\$599.99](#)




Moore's Law



Microprocessor Transistor Counts 1971-2011 & Moore's Law



Moore's Law and Performance



Parallelism is here... And Growing!



Parallelism for the Masses
"Opportunities and Challenges"

© Intel Corporation



3

64 core blade servers (\$6K) (shared memory)



AMD Opteron (sixteen-core) Model 6274

by [AMD](#)

(1 customer review)

List Price: \$693.00

Price: **\$599.99** ✓[Prime](#)

You Save: \$93.01 (13%)

Only 1 left in stock (more on the way).

Ships from and sold by [Amazon.com](#). Gift-wrap available.

Want it delivered Monday, November 5? Order it in the next 14 hours and 37 minutes

Delivery may be impacted by Hurricane Sandy. Proceed to checkout to see estimated

[43 new](#) from \$599.99

× 4 =



1024 “cuda” cores

amazon.com

Hello. [Sign in](#) to get personalized recommendations. New customer? [Start here.](#)

Your Amazon.com | Today's Deals | [Gifts & Wish Lists](#) | [Gift Cards](#)

Shop All Departments Search Electronics

All Electronics Brands Best Sellers Audio & Home Theater Camera & Photo Car E



EVGA GeForce GTX 590 Classified :
3DVI/Mini-Display Port SLI Ready Li
03G-P3-1596-AR

by [EVGA](#)

(16 customer reviews) | Like (29)

Price: \$924.56

In Stock.

Ships from and sold by [J-Electronics](#).

Only 1 left in stock--order soon.

[5 new](#) from \$749.99 [2 used](#) from \$695.00

Samsung Galaxy S IV to feature Exynos 28nm quad-core processor?

Written by [Andre Yoskowitz](#) @ 01 Nov 2012 18:02



It has been a few weeks but there is a new rumor regarding the upcoming [Samsung Galaxy S IV](#).

According to [reports](#), Samsung will pack next year's flagship device with its "Adonis" Exynos processor, a quad-core [ARM 15](#) beast that uses efficient 28nm tech.

Samsung is supposedly still testing the [application](#) processor, but mass production is scheduled for the Q1 2013 barring any delays.

Circa November 2012

Samsung Galaxy S IV is now Official: Octa-Core CPU, 5" Full HD Display & 13MP Camera

Follow: [Phones](#) [GT-I9500](#) [Samsung Display](#) [Samsung Exynos](#) [Samsung Galaxy S IV](#) [Samsung Mobile Unpacked 2013](#)



Samsung has just announced the Samsung Galaxy S4 at their Mobile Unpacked Event 2013 Episode 1 in New York, USA. The Galaxy S4 features a stunning 4.99" Full HD (1920×1080) SuperAMOLED display. With a 441 ppi pixel density, your eyes won't be able to distinguish the pixels, which ensures excellent visual comfort. Even though the Galaxy S4 has a large display and a massive battery of 2,600 MAh, it's only 7.9mm thick. Samsung's latest

flagship device is PACKED with powerful components, consisting of Samsung's latest Exynos 5 Octa-Core (5410) CPU based on ARM's big.LITTLE technology with Quad Cortex-

Intel Has a 48-Core Chip for Smartphones and Tablets

By Wolfgang Gruener OCTOBER 31, 2012 9:20 AM - Source: Computerworld

Intel has developed a prototype of a 48-core processor for smartphones. Before you ask: No, you can't buy a 48-core smartphone next year.



Parallel Hardware

- Many forms of parallelism
 - Supercomputers: large scale, shared memory
 - Clusters and data centers: large-scale, distributed memory
 - Multicores: tightly coupled, smaller scale
- Parallelism is important in the real world.

Key Challenge: Software (How to Write Parallel Code?)

- At a high-level, it is a two step process:
 - Design a work-efficient, low-span parallel algorithm
 - Implement it on the target hardware
- In reality: each system required different code because programming systems are immature
 - Huge effort to generate efficient parallel code.
 - Example: Quicksort in MPI is 1700 lines of code, and about the same in CUDA
 - Implement one parallel algorithm: a whole thesis.

Teaching Challenge: How to Teach Parallelism

- Education must prepare for the “real world” and research.
- Essential to teach parallel algorithm design.
- But many different models of parallelism.
- Inadequate parallel systems make practical work gruesome. Practically impractical in the context of an undergraduate course.

15-210 Approach

Enable parallel thinking by raising abstraction level


- I. **Parallel thinking:** Applicable to many machine models and programming languages
- II. **Reason** about correctness and efficiency of algorithms and data structures.

Parallel Thinking

- Recognizing true dependences: unteach sequential programming.
- Parallel algorithm-design techniques
 - Operations on aggregates: map/reduce/scan
 - Divide & conquer, contraction
 - Viewing computation as DAG
- Cost model based on work and span


Quicksort from Aho-Hopcroft-Ullman (1974)

```
procedure QUICKSORT(S):
    if S contains at most one element then return S
    else
        begin
            choose an element a randomly from S;
            let  $S_1$ ,  $S_2$  and  $S_3$  be the sequences of
                elements in S less than, equal to,
                and greater than a, respectively;
            return (QUICKSORT( $S_1$ ) followed by  $S_2$ 
                    followed by QUICKSORT( $S_3$ ))
    end
```



Quicksort from Sedgewick (2003)

```
public void quickSort(int[] a, int left, int right) {  
    int i = left-1;  int j = right;  
    if (right <= left) return;  
    while (true) {  
        while (a[++i] < a[right]);  
        while (a[right] < a[--j])  
            if (j==left) break;  
        if (i >= j) break;  
        swap(a,i,j); }  
        swap(a, i, right);  
        quickSort(a, left, i - 1);  
        quickSort(a, i+1, right); }
```



Programming Parallel Algorithms

NESL

Cilk

Scala

Java

OpenMP

F#


MPI

Parallel ML

C#

Threading Building Blocks (TBB)


OpenMP Overview



Slides by Matthew J. Sottile, Timothy G. Mattson, and Craig E Rasmussen

OpenMP Execution Model

- Fork-join parallelism




- Traditionally flat parallelism
- Recently nested parallelism. But not very good.

Example: Numerical Integration

Mathematically, we know that:

$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$



We can approximate the integral as a sum of rectangles:

$$\sum_{i=0}^N F(x_i) \Delta x \approx \pi$$

where each rectangle has width Δx and height $F(x_i)$ at the middle of interval i .

The code for Approximating PI

```
static long num_steps = 100000;
double step;
void main ()
{
    int i;  double x, pi, sum = 0.0;

    step = 1.0/(double) num_steps;
    x = 0.5 * step;
    for (i=0;i<= num_steps; i++){
        x+=step;
        sum += 4.0/(1.0+x*x);
    }
    pi = step * sum;
}
```

The code for Approximating PI

```
#include <omp.h>
static long num_steps = 100000;      double step;
void main ()
{   int i;           double x, pi, sum = 0.0;
    step = 1.0/(double) num_steps;
#pragma omp parallel for private(i, x) reduction(+:sum)
    for (i=0;i<= num_steps; i++){
        x = (i+0.5)*step;
        sum = sum + 4.0/(1.0+x*x);
    }
    pi = step * sum;
}
```

Private clause
creates data local to
a thread

Reduction used to
manage
dependencies

Cilk/Cilk++: extending C/C++ with fork-join parallelism for shared-memory systems

Slides by Charles E. Leiserson

Supercomputing Technologies Research Group
Computer Science and Artificial Intelligence Laboratory
Massachusetts Institute of Technology

Cilk Is Simple

- Cilk extends the C language with just a *handful* of keywords.
- Every Cilk program has a *serial semantics*.
- Not only is Cilk fast, it provides *performance guarantees* based on performance abstractions.
- Cilk is *processor-oblivious*.
- Cilk's *provably good* runtime system automatically manages low-level aspects of parallel execution, including protocols, load balancing, and scheduling.
- Cilk supports *speculative* parallelism.

An Example: Fibonacci

```
int fib (int n) {  
    if (n<2) return (n);  
    else {  
        int x,y;  
        x = fib(n-1);  
        y = fib(n-2);  
        return (x+y);  
    }  
}
```

C elision

Cilk code

```
cilk int fib (int n) {  
    if (n<2) return (n);  
    else {  
        int x,y;  
        x = spawn fib(n-1);  
        y = spawn fib(n-2);  
        sync;  
        return (x+y);  
    }  
}
```

Cilk is a ***faithful*** extension of C. A Cilk program's ***serial elision*** is always a legal implementation of Cilk semantics. Cilk provides ***no*** new data types.

Basic Cilk Keywords

```
cilk int fib (int n) {  
    if (n<2) return (n);  
    else {  
        int x,y;  
        x = spawn fib(n-1);  
        y = spawn fib(n-2);  
        sync;  
        return (x+y);  
    }  
}
```

Identifies a function as a **Cilk procedure**, capable of being spawned in parallel.

The named **child** Cilk procedure can execute in parallel with the **parent** caller.


Control cannot pass this point until all spawned children have returned.

Dynamic Multithreading

```
cilk int fib (int n) {  
    if (n<2) return (n);  
    else {  
        int x,y;  
        x = spawn fib(n-1);  
        y = spawn fib(n-2);  
        sync;  
        return (x+y);  
    }  
}
```

Example: **fib(4)**

Processor oblivious



The ***computation dag***
unfolds dynamically.

Example: Parallelizing Vector Addition

C

```
void vadd (real *A, real *B, int n){  
    int i; for (i=0; i<n; i++) A[i]+=B[i];  
}
```

Parallelizing Vector Addition

C

```
void vadd (real *A, real *B, int n){  
    int i; for (i=0; i<n; i++) A[i]+=B[i];  
}
```

C

```
void vadd (real *A, real *B, int n){  
    if (n<=BASE) {  
        int i; for (i=0; i<n; i++) A[i]+=B[i];  
    } else {  
        vadd (A, B, n/2);  
        vadd (A+n/2, B+n/2, n-n/2);  
    }  
}
```

Parallelization strategy:

1. Convert loops to recursion.

Parallelizing Vector Addition

C

```
void vadd (real *A, real *B, int n){  
    int i; for (i=0; i<n; i++) A[i]+=B[i];  
}
```

Cilk


```
cilk vadd (real *A, real *B, int n){  
    if (n<=BASE) {  
        int i; for (i=0; i<n; i++) A[i]+=B[i];  
    } else {  
        vaddn(A, B, n/2);  
        vaddn(A+n/2, B+n/2, n-n/2);  
    } sync;  
}
```

Parallelization strategy:

1. Convert loops to recursion.
2. Insert Cilk keywords.

Side benefit:
D&C is generally
good for caches!

How do the problems do on a modern multicore




Demo Time

Concurrency versus Parallelism

- ☞ **Concurrency:** non-determinacy due to interleaving multiple processes. Property of the application.
- ☞ **Parallelism:** using multiple processors/cores running at the same time. Property of the machine.

		Concurrency	
		Deterministic	Concurrent
Parallelism	Serial	Traditional programming	Traditional OS
	Parallel	Deterministic parallelism	Concurrent parallelism



Concurrency is Hard

Stack Example



- Serial
- Using locks
- Using compare and swap
- Using transactions



This is just a simple stack.

Concurrency : Stack Example 1

```
struct link {int v; link* next;}  
  
struct stack {  
    link* headPtr;  
  
    void push(link* a) {  
        a->next = headPtr;  
        headPtr = a;      }  
  
    link* pop() {  
        link* h = headPtr;  
        if (headPtr != NULL)  
            headPtr = headPtr->next;  
        return h; }  
}
```



Concurrency : Stack Example 1


```
struct link {int v; link* next;}

struct stack {
    link* headPtr;

    void push(link* a) {
        a->next = headPtr;
        headPtr = a;      }

    link* pop() {
        link* h = headPtr;
        if (headPtr != NULL)
            headPtr = headPtr->next;
        return h; }

}
```



Concurrency : Stack Example 1


```
struct link {int v; link* next;}

struct stack {
    link* headPtr;

    void push(link* a) {
        a->next = headPtr;
        headPtr = a;      }

    link* pop() {
        link* h = headPtr;
        if (headPtr != NULL)
            headPtr = headPtr->next;
        return h; }

}
```



Concurrency : Stack Example 1


```
struct link {int v; link* next;}

struct stack {
    link* headPtr;

    void push(link* a) {
        a->next = headPtr;
        headPtr = a;      }

    link* pop() {
        link* h = headPtr;
        if (headPtr != NULL)
            headPtr = headPtr->next;
        return h; }

}
```



Stack Example

- Serial
- Using locks
- Using compare and swap
- Using transactions



This is just a simple stack.

Concurrency : Stack Example 2

```
struct link {int v; link* next;}

struct stack {
    link* headPtr;

    void push(link* a) {
        withlock {
            a->next = headPtr;
            headPtr = a;    } }

    link* pop() {
        withlock {
            link* h = headPtr;
            if (headPtr != NULL)
                headPtr = headPtr->next;
            return h; } }

}
```

Concurrency : Stack Example 2

Locks are bad for many reasons

- Easy to create deadlocks
- If a process stalls in a lock it holds everyone else up

Stack Example

- Serial
- Using locks
- Using compare and swap
- Using transactions



This is just a simple stack.

Concurrency : CAS


The Compare and Swap instruction is a built in instruction on many machines that atomically Compares a value in a register to one in memory and swaps in a new value if equal

```
bool CAS(int *ptr, int oldv, in newv) {  
    if (*ptr = oldv) {  
        *ptr = newv;  
        return true;}  
    } else return false;  
}
```

On x86 instruction set implemented with
lock cmpxchg ptr,newv
oldv in EAX, puts flag in ZF


Concurrency : Stack Example 3

```
struct stack {  
    link* headPtr;  
  
    void push(link* a) {  
        do {  
            link* h = headPtr;  
            a->next = h;  
            while (!CAS(&headPtr, h, a)); }  
  
    link* pop() {  
        do {  
            link* h = headPtr;  
            if (h == NULL) return NULL;  
            link* nxt = h->next;  
            while (!CAS(&headPtr, h, nxt)) }  
        return h; }  
    }
```




Concurrency : Stack Example 3

```
struct stack {  
    link* headPtr;  
  
    void push(link* a) {  
        do {  
            link* h = headPtr;  
            a->next = h;  
            while (!CAS(&headPtr, h, a)); }  
  
    link* pop() {  
        do {  
            link* h = headPtr;  
            if (h == NULL) return NULL;  
            link* nxt = h->next;  
            while (!CAS(&headPtr, h, nxt)) }  
        return h; }  
    }
```




Concurrency : Stack Example 3

```
struct stack {  
    link* headPtr;  
  
    void push(link* a) {  
        do {  
            link* h = headPtr;  
            → a->next = h;  
            while (!CAS(&headPtr, h, a)); }  
  
    link* pop() {  
        do {  
            link* h = headPtr;  
            if (h == NULL) return NULL;  
            link* nxt = h->next;  
            while (!CAS(&headPtr, h, nxt)) }  
        return h; }  
    }
```



Concurrency : Stack Example 3

```
struct stack {  
    link* headPtr;  
  
    void push(link* a) {  
        do {  
            link* h = headPtr;  
            a->next = h;  
        }  
        while (!CAS(&headPtr, h, a)); }  
  
    link* pop() {  
        do {  
            link* h = headPtr;  
            if (h == NULL) return NULL;  
            link* nxt = h->next;  
        }  
        while (!CAS(&headPtr, h, nxt)); }  
        return h; }  
}
```



Concurrency : Stack Example 3'

```
P1 : x = s.pop();    y = s.pop();    s.push(x);
```

```
P2 : z = s.pop();
```



P2: h = headPtr;
P2: nxt = h->next;
P1: everything
P2: CAS (&headPtr, h, nxt)

The ABA problem

Can be fixed with counter and 2CAS, but...

Stack Example

- Serial
- Using locks
- Using compare and swap
- Using transactions

This is just a simple stack.



Concurrency : Stack Example 4

```
struct link {int v; link* next;}

struct stack {
    link* headPtr;           "using transactional memory"
    void push(link* a) {
        atomic {
            a->next = headPtr;
            headPtr = a;    } }

    link* pop() {
        atomic {
            link* h = headPtr;
            if (headPtr != NULL)
                headPtr = headPtr->next;
            return h; } }

}
```

Concurrency : Stack Example 3'

```
void swapTop(stack s) {  
    link* x = s.pop();  
    link* y = s.pop();  
    push(x);  
    push(y);  
}
```

Queues are trickier than stacks.

Complications due to **non-determinacy** in interleaving of instructions/operations on shared data

Summary

- Hardware parallelism
 - Plenty and ubiquitous: datacenters, mainframes, GPUs, multicore chips on desktops and mobiles
- Parallel software is a major challenge
 - Good solutions in specific domains.
 - Inadequate programming systems.
- 15210
 - Parallel thinking at a high abstraction level
 - Parallelism, not concurrency. Avoid concurrency if not needed (e.g., for deterministic computations)